

Phase 1 Project: Data Science

Project Details

- **Start Date:** 1 November 2023
- **Deadline:** 15 November 2023

Project Steps

1. Accessing Public API and Creating a Dataset

Used Python with requests library to access a public API (e.g., OpenWeatherMap API) and saved the obtained data as a CSV file.

```
In [1]: import pandas as pd
import requests

# URL for JSONPlaceholder API to fetch dummy data (example: posts)
api_url = 'https://jsonplaceholder.typicode.com/posts'

# Make a GET request to the API
response = requests.get(api_url)

if response.status_code == 200:
    data = response.json()

    # Creating a Pandas DataFrame from the retrieved data
    df = pd.DataFrame(data)

    # Saving the DataFrame to a CSV file
    df.to_csv('data.csv', index=False)
    print("CSV file 'data.csv' has been created with the fetched data.")
else:
    print("Failed to retrieve data. Status code:", response.status_code)
```

CSV file 'data.csv' has been created with the fetched data.

```
In [2]: print(df.head())
```

	userId	id	title \	body
0	1	1	sunt aut facere repellat provident occaecati e...	quia et suscipit\nsuscipit recusandae consequu...
1	1	2	qui est esse	est rerum tempore vitae\nsequi sint nihil repr...
2	1	3	ea molestias quasi exercitationem repellat qui...	et iusto sed quo iure\nvoluptatem occaecati om...
3	1	4	eum et est occaecati	ullam et saepe reiciendis voluptatem adipisci\...
4	1	5	nesciunt quas odio	repudiandae veniam quaerat sunt sed\nalias aut...

Data Cleaning with Pandas

This Python code snippet demonstrates the process of cleaning a dataset using Pandas, focusing on handling missing values and removing outliers.

Importing Necessary Libraries

The code begins by importing the required Python libraries, mainly `pandas` for data manipulation and handling.

Loading the Dataset

It retrieves the dataset from the specified URL using `pd.read_csv()` and stores it in a Pandas DataFrame.

Displaying Basic Information

It prints basic information about the dataset using `df.info()`, which includes column names, data types, and missing values.

Handling Missing Values

Missing values in numerical columns are filled with the mean of the respective column using `df.fillna(df.mean(), inplace=True)`. This replaces NaN values with the mean.

Removing Outliers

A common approach to removing outliers is applied. The code uses a threshold (in this case, 3 standard deviations from the mean) to filter out rows where a specific column's values are considered outliers.

The mean and standard deviation of the column are calculated, and then rows where the column value is beyond the threshold are filtered using boolean indexing.

Displaying the Cleaned Dataset

The head of the cleaned dataset is printed to display the first few rows after handling missing values and removing outliers.

Saving the Cleaned Dataset

Finally, the cleaned dataset is saved to a new CSV file named `cleaned_data.csv` using `df.to_csv()`.

Please note: The code is a basic example. Adjustments are necessary based on the characteristics of the dataset and specific requirements for handling missing values and outliers.

```
In [3]: import pandas as pd

# Load the downloaded dataset
file_path = 'dataset - netflix1.csv' # Replace with the actual file path
```

```
df = pd.read_csv(file_path)

# Display basic information about the dataset
print(df.info())

# Check for missing values
print(df.isnull().sum())

# Handle missing values (example: replacing missing values in 'director' column with 'Unknown')
df['director'].fillna('Unknown', inplace=True)

# Convert date columns to datetime format
df['date_added'] = pd.to_datetime(df['date_added'])
df['release_year'] = pd.to_datetime(df['release_year'], format='%Y') # Adjust format if needed

# Clean up 'duration' column (example: extract numeric values)
df['duration'] = df['duration'].str.extract('(\d+)').astype(float) # Extracting numeric values

# Display the head of the cleaned dataset
print(df.head())

# Save the cleaned dataset to a new CSV file
df.to_csv('cleaned_data.csv', index=False)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8790 entries, 0 to 8789
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   show_id               8790 non-null   object
1   type                  8790 non-null   object
2   title                 8790 non-null   object
3   director              8790 non-null   object
4   country               8790 non-null   object
5   date_added            8790 non-null   object
6   release_year          8790 non-null   int64
7   rating                8790 non-null   object
8   duration              8790 non-null   object
9   listed_in            8790 non-null   object
dtypes: int64(1), object(9)
memory usage: 686.8+ KB
None
show_id      0
type          0
title        0
director      0
country       0
date_added    0
release_year  0
rating        0
duration      0
listed_in     0
dtype: int64
show_id  type          title          director \
0      s1    Movie      Dick Johnson Is Dead  Kirsten Johnson
1      s3  TV Show      Ganglands      Julien Leclercq
2      s6  TV Show      Midnight Mass    Mike Flanagan
3     s14    Movie  Confessions of an Invisible Girl  Bruno Garotti
4      s8    Movie      Sankofa      Haile Gerima

country date_added release_year rating  duration \
0  United States 2021-09-25  2020-01-01  PG-13      90.0
1      France 2021-09-24  2021-01-01  TV-MA       1.0
2  United States 2021-09-24  2021-01-01  TV-MA       1.0
3      Brazil 2021-09-22  2021-01-01  TV-PG      91.0
4  United States 2021-09-24  1993-01-01  TV-MA     125.0

listed_in
0      Documentaries
1  Crime TV Shows, International TV Shows, TV Act...
2      TV Dramas, TV Horror, TV Mysteries
3      Children & Family Movies, Comedies
4  Dramas, Independent Movies, International Movies
```

```
In [4]: # Display basic information about the dataset
print(df.info())

# Identify and remove outliers from the 'duration' column
column_name = 'duration' # Replace with the column name containing outliers

# Calculate mean and standard deviation
mean = df[column_name].mean()
std_dev = df[column_name].std()

# Define a threshold (e.g., 3 standard deviations from the mean)
threshold = 3 * std_dev

# Filter out rows where the column value is considered an outlier
```

```
df = df[abs(df[column_name] - mean) < threshold]

# Display the head of the cleaned dataset after outlier removal
print(df.head())

# Save the cleaned dataset to a new CSV file
df.to_csv('cleaned_data_without_outliers.csv', index=False)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8790 entries, 0 to 8789
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   show_id               8790 non-null   object
1   type                  8790 non-null   object
2   title                 8790 non-null   object
3   director              8790 non-null   object
4   country               8790 non-null   object
5   date_added            8790 non-null   datetime64[ns]
6   release_year          8790 non-null   datetime64[ns]
7   rating                8790 non-null   object
8   duration              8790 non-null   float64
9   listed_in             8790 non-null   object
dtypes: datetime64[ns](2), float64(1), object(7)
memory usage: 686.8+ KB
None
```

	show_id	type	title	director
0	s1	Movie	Dick Johnson Is Dead	Kirsten Johnson
1	s3	TV Show	Ganglands	Julien Leclercq
2	s6	TV Show	Midnight Mass	Mike Flanagan
3	s14	Movie	Confessions of an Invisible Girl	Bruno Garotti
4	s8	Movie	Sankofa	Haile Gerima

	country	date_added	release_year	rating	duration
0	United States	2021-09-25	2020-01-01	PG-13	90.0
1	France	2021-09-24	2021-01-01	TV-MA	1.0
2	United States	2021-09-24	2021-01-01	TV-MA	1.0
3	Brazil	2021-09-22	2021-01-01	TV-PG	91.0
4	United States	2021-09-24	1993-01-01	TV-MA	125.0

	listed_in
0	Documentaries
1	Crime TV Shows, International TV Shows, TV Act...
2	TV Dramas, TV Horror, TV Mysteries
3	Children & Family Movies, Comedies
4	Dramas, Independent Movies, International Movies

```
In [5]: # outliers as been removed(8790-8781=9)
print(df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 8781 entries, 0 to 8789
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   show_id                8781 non-null   object
1   type                   8781 non-null   object
2   title                  8781 non-null   object
3   director               8781 non-null   object
4   country                8781 non-null   object
5   date_added             8781 non-null   datetime64[ns]
6   release_year           8781 non-null   datetime64[ns]
7   rating                 8781 non-null   object
8   duration               8781 non-null   float64
9   listed_in              8781 non-null   object
dtypes: datetime64[ns](2), float64(1), object(7)
memory usage: 754.6+ KB
None

```

Phase -2

Analysis of Dataset and Graph Generation

Project Details

- **Start Date:** 16 November 2023
- **Deadline:** 1 December 2023
- **Task:** Analyzing the provided dataset and creating graphs using Seaborn and Matplotlib.

Dataset

- The dataset used for analysis can be accessed [here](#).

Methodology

1. Loading the Dataset:

- The dataset was imported into a Pandas DataFrame for analysis using Python.
- Python libraries used: Pandas, Seaborn, Matplotlib.

2. Data Exploration:

- Performed data exploration to understand the structure, columns, and types of data present in the dataset.
- Utilized commands like `df.head()`, `df.info()`, and `df.describe()` to gain insights into the data.

3. Data Visualization:

- Utilized Seaborn and Matplotlib to generate visualizations and insights from the dataset.
- Created various types of plots such as:

- Histograms
- Scatter Plots
- Bar Charts
- Pair Plots

4. Findings and Analysis:

- Interpreted the visualizations and derived insights from the generated graphs.
- Noted any trends, patterns, or anomalies observed in the data.

Code Snippets

Example code used for generating graphs:

```
```python import pandas as pd import seaborn as sns import matplotlib.pyplot as plt
```

## Load the dataset

```
df = pd.read_csv('your_file_path.csv')
```

## Creating a histogram

```
sns.histplot(data=df, x='column_name', kde=True) plt.title('Histogram of column_name')
plt.show()
```

```
In [6]: import seaborn as sns
import matplotlib.pyplot as plt


print(df.info())

Example visualizations
Plotting count of 'type' (Movie vs TV Show)
sns.countplot(x='type', data=df)
plt.title('Count of Movies and TV Shows')
plt.xlabel('Type')
plt.ylabel('Count')
plt.show()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 8781 entries, 0 to 8789
Data columns (total 10 columns):
 # Column Non-Null Count Dtype
--- ---
 0 show_id 8781 non-null object
 1 type 8781 non-null object
 2 title 8781 non-null object
 3 director 8781 non-null object
 4 country 8781 non-null object
 5 date_added 8781 non-null datetime64[ns]
 6 release_year 8781 non-null datetime64[ns]
 7 rating 8781 non-null object
 8 duration 8781 non-null float64
 9 listed_in 8781 non-null object
dtypes: datetime64[ns](2), float64(1), object(7)
memory usage: 754.6+ KB
None

```




```

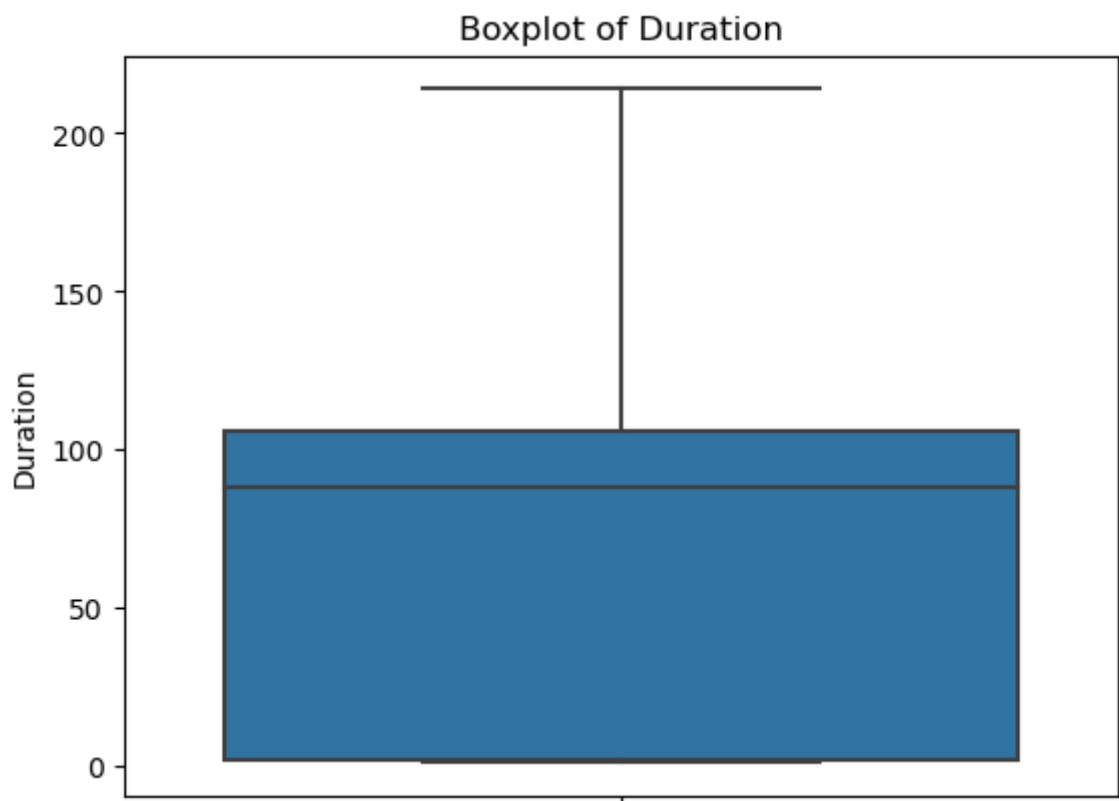
In [7]: # Creating a histogram of 'release_year'
sns.histplot(data=df, x='release_year', bins=20, kde=True)
plt.title('Distribution of Release Years')
plt.xlabel('Release Year')
plt.ylabel('Frequency')
plt.show()

```





```
In [8]: # Create a boxplot for 'duration' to detect outliers
sns.boxplot(data=df, y='duration')
plt.title('Boxplot of Duration')
plt.ylabel('Duration')
plt.show()
```



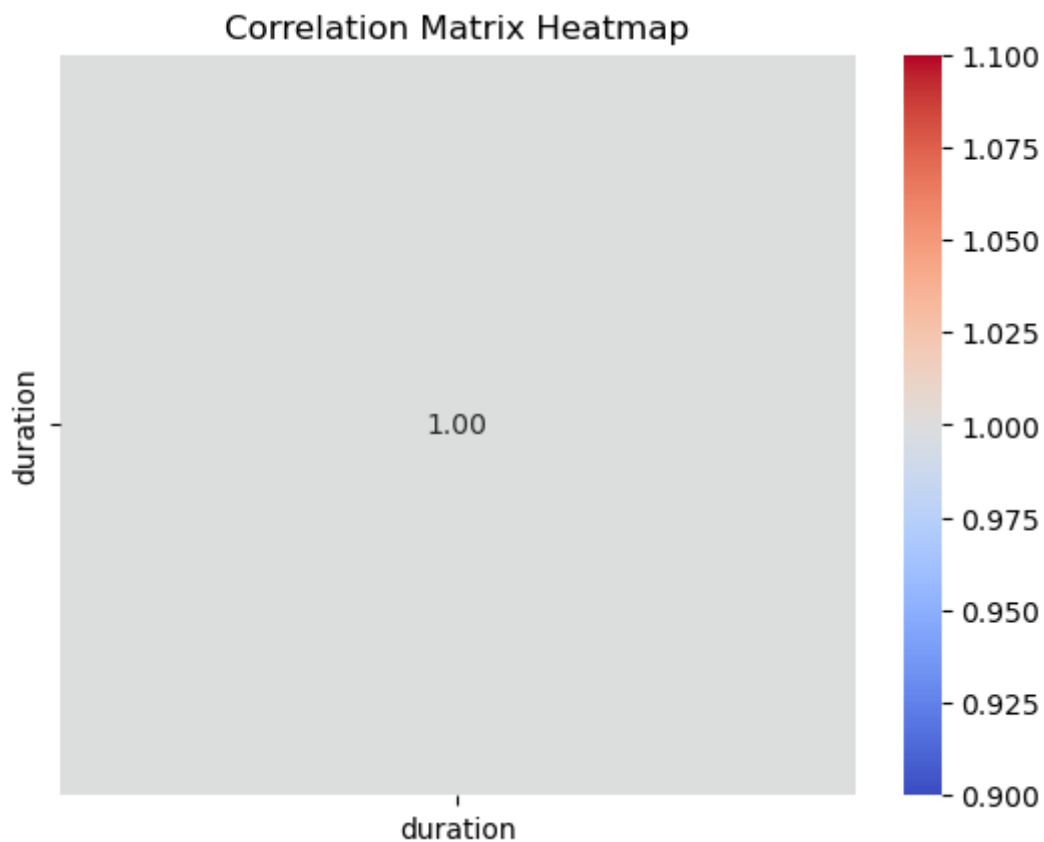
```
In [9]: # Drop non-numeric or irrelevant columns for correlation
non_numeric_cols = ['show_id'] # Adjust with other non-numeric columns if needed
```

```
df_numeric = df.select_dtypes(include=['int64', 'float64'])

Display basic information about the remaining numeric columns
print(df_numeric.info())

Correlation matrix heatmap for numeric columns
sns.heatmap(df_numeric.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix Heatmap')
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 8781 entries, 0 to 8789
Data columns (total 1 columns):
 # Column Non-Null Count Dtype
--- -
 0 duration 8781 non-null float64
dtypes: float64(1)
memory usage: 137.2 KB
None
```



## Linear Regression Using Provided Datasets

### Objective

The goal is to demonstrate the process of training a simple linear regression model using the provided datasets. The training dataset consists of 'x' and 'y' values, and the test dataset contains 'x' values for which predictions need to be made.

### Steps

#### Step 1: Load the Training Data

- Utilize the Pandas library to fetch data from the provided Google Sheets URLs for the training dataset.
- Display the initial rows to understand the data structure.

```
python import pandas as pd
```

```
train_data_url = 'Training_Data_Google_Sheets_URL' train_data = pd.read_csv(train_data_url)
print(train_data.head())
```

```
In [10]: import pandas as pd

Load the training dataset
train_data_url = 'https://docs.google.com/spreadsheets/d/e/2PACX-1vRTK2NvcndgPX41C:
train_data = pd.read_csv(train_data_url)

Display the first few rows of the training data
print(train_data.head())
```

	x	y
0	24.0	21.549452
1	50.0	47.464463
2	15.0	17.218656
3	38.0	36.586398
4	87.0	87.288984

```
In [14]: # Check for NaN values in the 'y' column of the training data
nan_values = train_data['y'].isnull().sum()
print("Number of NaN values in 'y' column:", nan_values)

train_data = train_data.dropna(subset=['y'])

train_data['y'].fillna(train_data['y'].mean(), inplace=True)

Number of NaN values in 'y' column: 0
```

## Step 2: Train a Simple Linear Regression Model

Use a machine learning library like Scikit-learn to create and train a linear regression model.

```
In [15]: from sklearn.linear_model import LinearRegression

Create the linear regression model
model = LinearRegression()

Fit the model using the training data
X_train = train_data[['x']] # Features
y_train = train_data['y'] # Target variable
model.fit(X_train, y_train)
```

```
Out[15]: ▾ LinearRegression
LinearRegression()
```

## Step 3: Load the Test Data and Make Predictions

Repeat a similar process for the test dataset.

```
In [16]: # Load the test dataset
test_data_url = 'https://docs.google.com/spreadsheets/d/e/2PACX-1vRyvZ71knwiSghK9a6
test_data = pd.read_csv(test_data_url)

Display the first few rows of the test data
print(test_data.head())

Perform predictions using the trained model
X_test = test_data[['x']] # Features
predictions = model.predict(X_test)

Display the predictions
print(predictions)
```

	x	y					
0	77	79.775152					
1	21	23.177279					
2	22	25.609262					
3	20	17.857388					
4	36	41.849864					
[	76.94327594	20.90651855	21.90717494	19.90586217	35.91636428	14.90258026	
	61.93343021	94.95509081	19.90586217	4.89601644	3.89536006	18.90520579	
	95.95574719	61.93343021	35.91636428	14.90258026	64.93539936	13.90192388	
	86.94983976	68.93802488	88.95115252	50.92621001	88.95115252	26.91045685	
	96.95640358	57.93080468	78.9445887	20.90651855	92.95377805	26.91045685	
	98.95771634	30.91308237	32.91439514	79.94524508	27.91111323	46.92358448	
	52.92752277	68.93802488	27.91111323	32.91439514	90.95246528	70.93933765	
	49.92555363	75.94261956	3.89536006	36.91702066	69.93868127	67.9373685	
	39.91898981	34.9157079	93.95443443	87.95049614	51.92686639	30.91308237	
	58.93146107	-0.10726546	38.91833343	63.93474297	68.93802488	56.9301483	
	12.9012675	71.93999403	75.94261956	60.93277383	81.94655785	17.90454941	
	40.91964619	49.92555363	54.92883554	12.9012675	45.9229281	12.9012675	
	78.9445887	52.92752277	14.90258026	27.91111323	80.94590147	68.93802488	
	51.92686639	83.94787061	67.9373685	26.91045685	55.92949192	47.92424086	
	39.91898981	38.91833343	81.94655785	99.95837272	58.93146107	42.92095896	
	66.93671212	37.91767705	62.93408659	90.95246528	59.93211745	13.90192388	
	20.90651855	86.94983976	72.94065041	31.91373876	1.8940473	81.94655785	
	18.90520579	73.94130679	41.92030257	11.90061112	0.89339092	89.9518089	
	88.95115252	-0.10726546	40.91964619	15.90323665	93.95443443	96.95640358	
	65.93605574	23.9084877	16.90389303	89.9518089	12.9012675	-0.10726546	
	63.93474297	95.95574719	97.95705996	11.90061112	40.91964619	46.92358448	
	77.94393232	19.90586217	88.95115252	28.91176961	63.93474297	74.94196317	
	11.90061112	24.90914408	27.91111323	29.91242599	64.93539936	58.93146107	
	63.93474297	52.92752277	70.93933765	96.95640358	72.94065041	8.89864197	
	11.90061112	62.93408659	98.95771634	59.93211745	34.9157079	1.8940473	
	59.93211745	31.91373876	93.95443443	83.94787061	62.93408659	21.90717494	
	80.94590147	92.95377805	32.91439514	6.89732921	41.92030257	45.9229281	
	53.92817916	15.90323665	48.92489725	42.92095896	94.95509081	65.93605574	
	20.90651855	34.9157079	79.94524508	36.91702066	53.92817916	55.92949192	
	0.89339092	31.91373876	57.93080468	31.91373876	45.9229281	71.93999403	
	16.90389303	96.95640358	92.95377805	90.95246528	36.91702066	3.89536006	
	53.92817916	50.92621001	26.91045685	45.9229281	91.95312167	72.94065041	
	76.94327594	90.95246528	60.93277383	98.95771634	3.89536006	71.93999403	
	18.90520579	56.9301483	77.94393232	25.90980046	73.94130679	89.9518089	
	65.93605574	12.9012675	39.91898981	76.94327594	66.93671212	74.94196317	
	22.90783132	44.92227172	58.93146107	43.92161534	22.90783132	54.92883554	
	54.92883554	94.95509081	11.90061112	3.89536006	6.89732921	99.95837272	
	47.92424086	41.92030257	95.95574719	38.91833343	99.95837272	86.94983976	
	13.90192388	13.90192388	36.91702066	4.89601644	87.95049614	90.95246528	
	64.93539936	73.94130679	55.92949192	15.90323665	4.89601644	27.91111323	
	91.95312167	45.9229281	53.92817916	38.91833343	43.92161534	30.91308237	
	67.9373685	85.94918338	89.9518089	37.91767705	20.90651855	94.95509081	
	55.92949192	59.93211745	64.93539936	77.94393232	88.95115252	5.89667283	
	66.93671212	35.91636428	15.90323665	99.95837272	44.92227172	72.94065041	
	56.9301483	19.90586217	75.94261956	33.91505152	54.92883554	71.93999403	
	54.92883554	7.89798559	55.92949192	71.93999403	57.93080468	5.89667283	
	95.95574719	22.90783132	57.93080468	22.90783132	18.90520579	24.90914408	
	63.93474297	20.90651855	58.93146107	18.90520579	15.90323665	41.92030257	
	42.92095896	60.93277383	91.95312167	10.89995474	40.91964619	0.89339092	
	7.89798559	70.93933765	45.9229281	54.92883554	61.93343021	46.92358448]	

In [ ]: